

## A Formal Verification of Failover Impossibility

We formalize and mechanically verify three impossibility theorems for transparent RDMA failover. All proofs are verified in Rocq 9.0 and available at <https://github.com/taooceros/shift-verification>.

### A.1 Model and Definitions

**Definition** (Execution Trace). A trace  $\mathcal{T}$  is a sequence of events including: operation sends (`EvSend`), packet/ACK losses, executions at the receiver (`EvExecute`), completions, and timeouts.

**Definition** (Sender View). The *sender view*  $\sigma(\mathcal{T})$  projects a trace to only sender-observable events: sends, completions, and timeouts. Crucially, the sender cannot observe `EvExecute`, `EvPacketLost`, or `EvAckLost` directly.

**Definition** (Transparent Overlay). A failover mechanism is *transparent* if its retransmission decision  $D : \sigma(\mathcal{T}) \times \text{Op} \rightarrow \{\text{retransmit}, \text{skip}\}$  depends only on the sender view—no persistent metadata, no receiver-side modifications, no application protocol changes.

### A.2 Theorem 1: Indistinguishability of Packet Loss and ACK Loss

**Theorem** (Impossibility of Safe Retransmission). *For any transparent overlay  $D$ , there exist traces  $\mathcal{T}_1$  and  $\mathcal{T}_2$  such that  $\sigma(\mathcal{T}_1) = \sigma(\mathcal{T}_2)$ , but safety requires  $D(\sigma(\mathcal{T}_1)) = \text{retransmit}$  while  $D(\sigma(\mathcal{T}_2)) = \text{skip}$ .*

*Proof.* We construct two traces for a Write operation to address  $A_{\text{data}}$  with value  $V_1$ :

$\mathcal{T}_1$  (packet lost—retransmission required for liveness):

$$[\text{EvSend}(W), \text{EvPacketLost}(W), \text{EvTimeout}(W)]$$

$\mathcal{T}_2$  (ACK lost, memory reused—retransmission corrupts data):

$$[\text{EvSend}(W), \text{EvReceive}(W), \text{EvExecute}(W), \text{EvAppConsume}, \text{EvAppReuse}(V'), \text{EvAckLost}(W), \text{EvTimeout}(W)]$$

Both traces yield sender view  $\sigma(\mathcal{T}_1) = \sigma(\mathcal{T}_2) = [\text{ObsSent}(W), \text{ObsTimeout}(W)]$ .

- In  $\mathcal{T}_1$ : operation never executed  $\rightarrow$  liveness requires retransmission
- In  $\mathcal{T}_2$ : operation executed, receiver consumed  $V_1$  and wrote new value  $V'$   $\rightarrow$  retransmission would overwrite  $V'$  with stale  $V_1$

Since  $D$  is a function and  $\sigma(\mathcal{T}_1) = \sigma(\mathcal{T}_2)$ , we have  $D(\sigma(\mathcal{T}_1)) = D(\sigma(\mathcal{T}_2))$ . But the traces require opposite decisions. Contradiction.  $\square$

**Implication for SHIFT:** This theorem explains why SHIFT cannot guarantee correctness for *all* traffic. However, for protocols where the receiver does not access data until a subsequent signal (e.g., NCCL Simple’s flag mechanism), the “ACK lost + memory reused” scenario cannot occur before SHIFT completes retransmission.

### A.3 Theorem 2: Non-Idempotency of Operations

**Theorem** (FADD Non-Idempotency). *For any  $\delta > 0$  and memory  $m$ , FADD is not idempotent:  $\text{exec}_{\text{FADD}}(\text{exec}_{\text{FADD}(m, a, \delta)}, a, \delta) \neq \text{exec}_{\text{FADD}(m, a, \delta)}$ .*

**Theorem** (Queue Sliding (Two-Sided Non-Idempotency)). *Retrying a SEND operation is not idempotent because it consumes an additional Receive WQE.*

*Proof.* Let the receiver queue  $Q_R = [R_1, R_2, \dots]$ . Trace 1 (Success): Message  $M_1$  consumes  $R_1$ .  $Q_{R'} = [R_2, \dots]$ . ACK lost. Trace 2 (Retry): Message  $M_1$  (retry) consumes  $R_2$ .  $Q_{R''} = [R_3, \dots]$ . Result:  $M_1$  is duplicated in buffers  $B_1$  and  $B_2$ , and  $R_2$  (intended for  $M_2$ ) is lost. The streams are permanently misaligned.  $\square$

**Theorem** (CAS Double Success). *Under concurrent modification, a CAS retry can succeed even after the original succeeded.*

*Proof.* Consider sender  $S$  with  $\text{CAS}(a, 0, 1)$  and concurrent process  $P$  with  $\text{CAS}(a, 1, 0)$ :

Step	Actor	Operation	$m[a]$
0	—	Initial	0

1	$S$	CAS( $0 \rightarrow 1$ )	1 (success)
2	$S$	Fault before ACK	—
3	$P$	CAS( $1 \rightarrow 0$ )	0 (success)
4	$S$	Retry CAS( $0 \rightarrow 1$ )	1 (success!)

$S$ 's single logical CAS executed twice, and  $P$ 's modification was silently overwritten.  $\square$

#### A.4 Theorem 3: Consensus Hierarchy Barrier

We prove that correct failover requires solving 2-process consensus, which read-only verification cannot achieve.

##### A.4.1 Background: Herlihy's Consensus Hierarchy

**Definition** (Consensus Number).  $\text{CN}(X) = n$  means primitive  $X$  can solve wait-free  $n$ -process consensus but not  $(n+1)$ -consensus. Key results:  $\text{CN}(\text{Register}) = 1$ ,  $\text{CN}(\text{FADD}) = 2$ ,  $\text{CN}(\text{CAS}) = \infty$ .

The hierarchy is *strict*: primitives with  $\text{CN} = k$  cannot implement primitives with  $\text{CN} > k$ .

##### A.4.2 Reduction: Failover Solver $\Rightarrow$ 2-Consensus

**Definition** (Failover Solver).  $F : \text{Memory} \rightarrow \{\text{Commit}, \text{Abort}\}$  returns the correct decision: Commit if the original CAS executed, Abort otherwise.

**Theorem** (2-Consensus from Failover Solver). *A correct failover solver  $F$  yields a 2-consensus protocol:*

1. Each process  $P_i$  writes its input to  $\text{proposed}[i]$
2. Both call  $F(m)$  on the shared memory state
3. If  $F(m) = \text{Commit}$ : decide  $\text{proposed}[0]$
4. If  $F(m) = \text{Abort}$ : decide  $\text{proposed}[1]$

*Proof.*

- **Wait-free**: No loops; finite steps.
- **Agreement**: Both call same  $F$  on same  $m \rightarrow$  same result  $\rightarrow$  same decision.
- **Validity**: If CAS executed ( $P_0$  “won”),  $F(m) = \text{Commit}$ , decision =  $\text{proposed}[0]$ . Similarly for  $P_1$ .

$\square$

**Theorem** (Failover Solver Yields 2-Consensus). *A correct failover solver yields a read-based protocol solving 2-consensus.*

*Proof.* Given  $F$  satisfying `solves_failover`, construct observation  $\text{obs}(e, i) :=$  if  $F(m_0)$  then 0 else 1 (constant, trivially satisfies `valid_rw_observation`) and decision  $\text{decide}(x) :=$  if  $x = 0$  then 0 else 1.

- Solo  $P_0$ : `solves_failover` on `HistExecuted`( $m_0$ ) gives  $F(m_0) = \text{true} \rightarrow \text{obs} = 0$ ,  $\text{decide}(0) = 0 \checkmark$
- Solo  $P_1$ : `solves_failover` on `HistNotExecuted`( $m_0$ ) gives  $F(m_0) = \text{false} \rightarrow \text{obs} = 1$ ,  $\text{decide}(1) = 1 \checkmark$

$\square$

**Theorem** (Failover Impossible by Register CN=1). *No correct failover solver exists. The impossibility is formally derived from the Register CN=1 theorem.*

*Proof.* By mechanized reduction chain:

1. `failover_solver_yields_2consensus`: a correct solver yields `obs` (satisfying `valid_rw_observation`) and `decide` satisfying solo validity for both processes
2. `readwrite_2consensus_impossible_same_protocol`: no read-based `obs`/`decide` pair can satisfy both solo validities (Register CN=1)
3. `failover_impossible_by_read_cn`: combines (1) and (2) into contradiction

This is not a separate ABA argument—it is a formal *consequence* of  $\text{CN}(\text{Register}) = 1$ .  $\square$

**Theorem** (Main Impossibility). *Transparent failover for atomic operations is impossible.*

*Proof.* The main theorem (`transparent_cas_failover_impossible`) lifts the CN-based impossibility to the `TransparentFailover` interface:

1. Any `provides_reliable_cas` witness yields a `VerificationMechanism` satisfying `solves_failover`
2. `failover_impossible_by_read_cn` derives contradiction via the Register CN=1 barrier

□

## A.5 Mechanization Summary

Component	Lines	Key Theorems
Core (Memory, Ops, Traces)	400	<code>mem_read_write_same, exec_op</code>
Theorem 1	300	<code>sender_views_equal, impossibility_safe_retransmission</code>
Theorem 2	250	<code>fadd_not_idempotent, send_queue_sliding, cas_double_success</code>
Theorem 3	2200	<code>register_cn_1_verified, failover_impossible_by_read_cn, transparent_cas_failover_impossible</code>

Table 1: Rocq formalization statistics (total 3,900 lines)

All proofs compile with Rocq 9.0 without axioms beyond the standard library. The formalization models RDMA operations as state transformers ( $\text{Memory} \rightarrow \text{Memory} \times \text{Result}$ ), traces as event sequences, and the sender view as a projection function. The failover impossibility (Theorem 3) is formally derived FROM the Register CN=1 theorem via `failover_impossible_by_read_cn`, mechanizing the connection to Herlihy's hierarchy.

## A.6 Connection to SHIFT Design

These impossibility results directly inform SHIFT's design decisions:

Theorem	SHIFT Design Choice
Thm 1: Indistinguishability	Best-effort WR-level retransmission; error propagation when safety cannot be guaranteed
Thm 2: Non-idempotency (Atomics)	Return error if atomic WR is in-flight during fault
Thm 2: Queue Sliding (Two-Sided)	Implement 3-way handshake to re-synchronize queue indices (breaking transparency to ensure correctness)
Thm 3: Consensus barrier	No attempt to verify execution status via memory reads; rely on protocol-level idempotency or handshake instead

SHIFT's approach—supporting NCCL Simple while rejecting atomics and synchronizing two-sided ops—is not a limitation of implementation but a necessary consequence of these fundamental impossibility results. The boundary identified in Table 1 is precisely the boundary between what can and cannot be transparently failed over.