

Formal Verification of RDMA Failover Impossibility

We formalize and mechanically verify three impossibility theorems for transparent RDMA failover using the Rocq proof assistant (formerly Coq). All proofs are available at github.com/taooceros/shift-verification.

Theorem 1: Indistinguishability of Packet Loss and ACK Loss

Definition (Sender View). Let \mathcal{T} be an execution trace. The *sender view* $\sigma(\mathcal{T})$ is the projection containing only sender-observable events: operation sends, completions, and timeouts.

Definition (Transparent Overlay). A failover mechanism is *transparent* if its retransmission decision $D : \sigma(\mathcal{T}) \times \text{Op} \rightarrow \{0, 1\}$ depends only on the sender view.

Theorem (Impossibility of Safe Retransmission). For any transparent overlay D , there exist executions \mathcal{T}_1 (packet lost) and \mathcal{T}_2 (ACK lost, memory reused) such that:

$$\sigma(\mathcal{T}_1) = \sigma(\mathcal{T}_2) \quad (1)$$

but safety requires $D(\sigma(\mathcal{T}_1)) = 1$ (retransmit) while $D(\sigma(\mathcal{T}_2)) = 0$ (do not retransmit).

Proof. We construct two traces with identical sender views but opposite correctness requirements:

\mathcal{T}_1 : [Send(W_D), PacketLost(W_D), Timeout(W_D)]

\mathcal{T}_2 : [Send(W_D), Receive, Execute, AppConsume, AppReuse(V'), AckLost, Timeout(W_D)]

Both produce sender view [ObsSent(W_D), ObsTimeout(W_D)]. In \mathcal{T}_1 , the operation was never executed (liveness requires retry). In \mathcal{T}_2 , the operation executed and memory was reused with value $V' \neq V_1$ (safety forbids retry). Since D is a function, $D(\sigma(\mathcal{T}_1)) = D(\sigma(\mathcal{T}_2))$, contradicting the requirements. \square

Theorem 2: Non-Idempotency of Atomic Operations

Theorem (FADD Non-Idempotency). For any $\delta > 0$ and memory state m , FADD is not idempotent:

$$\text{exec}_{\text{FADD}}(\text{exec}_{\text{FADD}(m, a, \delta)}, a, \delta) \neq \text{exec}_{\text{FADD}(m, a, \delta)} \quad (2)$$

Proof. Let $m[a] = v$. After one FADD: $m'[a] = v + \delta$. After retry: $m''[a] = v + 2\delta$. Since $\delta > 0$, we have $v + \delta \neq v + 2\delta$. \square

Theorem (CAS Retry Violation). Under concurrent modification, a CAS retry can succeed twice, violating at-most-once semantics.

Proof. Consider sender S with $\text{CAS}(a, 0, 1)$ and concurrent process P with $\text{CAS}(a, 1, 0)$:

State 0: $m[a] = 0$

State 1: $S.\text{CAS}(0, 1)$ succeeds $\rightarrow m[a] = 1$

State 2: $P.\text{CAS}(1, 0)$ succeeds $\rightarrow m[a] = 0$

State 3: S retries $\text{CAS}(0, 1) \rightarrow$ succeeds again!

S 's single CAS executed twice, and P 's successful modification was silently overwritten. \square

Theorem 3: Consensus Hierarchy Barrier

We prove that failover coordination is equivalent to 2-process consensus, which read-only verification cannot solve.

Unified Observation Constraint Framework

Definition (Observation Constraint). Each synchronization primitive defines a constraint on what protocols can observe:

| Primitive | Constraint |
|-----------|--|
| Register | $\text{valid}_{\text{rw}} : \text{obs}(\text{exec}, i)$ depends only on writes before i |
| FADD | $\text{valid}_{\text{fadd}} : \text{obs}(\text{exec}, i)$ depends only on $\{j : j \text{ before } i\}$ (set, not order) |
| CAS | $\text{valid}_{\text{cas}} : \text{obs}(\text{exec}, i) = \text{winner}(\text{exec})$ (first process) |

The constraints are *derived* from primitive semantics:

- **Register**: Reads are invisible; only writes affect observable state
- **FADD**: Returns sum of prior deltas; $\delta_0 + \delta_1 = \delta_1 + \delta_0$
- **CAS**: First CAS to sentinel wins; all subsequent fail; all read same value

Consensus Number Verification

Definition (Consensus Number). $\text{CN}(X) = n$ iff X can solve n -consensus but not $(n+1)$ -consensus.

Lemma (Register CN = 1). For any observation function satisfying valid_{rw} , solo executions [0] and [1] produce identical observations (both have empty prior write history) but require decisions 0 and 1 respectively.

Lemma (FADD CN = 2). For any observation function satisfying $\text{valid}_{\text{fadd}}$, executions [0, 1, 2] and [1, 0, 2] are indistinguishable to process 2 (both see $\{0, 1\}$ ran before), but require decisions 0 and 1.

Lemma (CAS CN = ∞). Any observation function satisfying $\text{valid}_{\text{cas}}$ gives $\text{obs}(\text{exec}, i) = \text{winner}(\text{exec})$. Different winners \rightarrow different observations \rightarrow always distinguishable.

Reduction: Failover Solver \Rightarrow 2-Consensus Protocol

Following Herlihy's methodology (cf. Theorem 5.4.1 for FIFO queues), we prove failover requires $\text{CN} \geq 2$ by constructing a 2-consensus protocol from a hypothetical failover solver.

Definition (Failover Solver). A failover solver $F : \text{Memory} \rightarrow \{\text{true}, \text{false}\}$ returns **true** (Commit) if CAS was executed, **false** (Abort) otherwise.

Theorem (2-Consensus Protocol from Failover Solver). Given a correct failover solver F , the following protocol solves 2-consensus:

FIFO Protocol (Herlihy)

```

Queue := [WIN, LOSE]
proposed[i] := v_i
result := dequeue()
if result = WIN
    then decide(proposed[me])
else decide(proposed[other])

```

Failover Protocol (Ours)

```

Memory := m (ABA state)
proposed[i] := v_i
result := F(m)
if result = true
    then decide(proposed[0])
else decide(proposed[1])

```

The failover solver F acts like `dequeue()`: it reveals who “won”.

Proof. We verify the three conditions:

- **Wait-free:** The protocol contains no loops. Each process executes a finite number of steps. ✓
- **Agreement:** Both processes call the same F on the same memory m . They get the same result and select from the same `proposed[]` array. Therefore they decide the same value. ✓
- **Validity:** If P_0 won (CAS executed), then $F(m) = \text{true}$ by correctness of F , so both decide $\text{proposed}[0] = P_0$'s input. Similarly for P_1 . The decision is always the winner's input. ✓

□

Theorem (Failover Requires CN ≥ 2). No correct failover solver exists.

Proof. Suppose F is a correct failover solver. By correctness:

- $F(m) = \text{true}$ for $H_1 = \text{HistExecuted}(m)$
- $F(m) = \text{false}$ for $H_0 = \text{HistNotExecuted}(m)$

But both histories have the same final memory m (ABA problem). Thus $F(m)$ must equal both **true** and **false**. Contradiction.

Since a correct failover solver would yield a correct 2-consensus protocol, and no such solver exists, failover inherits the CN ≥ 2 requirement from 2-consensus. □

Theorem (Transparent Failover Impossibility). Read-only verification (the only tool available under transparency) has CN = 1 < 2. By Herlihy's impossibility theorem, CN = 1 primitives cannot solve 2-consensus. Therefore, transparent failover is impossible.

Theorem (Main Result). Transparent RDMA failover for atomic operations is impossible because:

1. Failover requires solving 2-consensus
2. Transparency limits verification to read-only operations
3. CN(Register) = 1 < 2
4. By Herlihy's hierarchy, CN=1 primitives cannot solve 2-consensus

Mechanization

| Component | Lines | Key Theorems |
|------------------|-------|--|
| Core definitions | 400 | Memory model, RDMA operations, traces |
| Theorem 1 | 200 | <code>impossibility_safe_retransmission</code> |
| Theorem 2 | 300 | <code>fadd_not_idempotent, cas_double_success</code> |
| Theorem 3 | 1200 | <code>register_cn_1_verified, fadd_cn_2_verified, valid_cas_no_ambiguity, transparent_cas_failover_impossible</code> |

Table 1: Rocq formalization statistics

All proofs are constructive and fully mechanized in Rocq 9.0. The consensus number framework provides a unified treatment where each primitive's limitation is derived from its operational semantics, and the failover impossibility follows as a direct consequence of Herlihy's hierarchy applied to the structural isomorphism between failover and 2-consensus.